

# VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



**LAB REPORT**  
**on**

## **Machine Learning (23CS6PCMAL)**

*Submitted by*

**Abhishek S Halagadagi (1BM22CS008)**

*in partial fulfillment for the award of the degree of*

**BACHELOR OF ENGINEERING**  
*in*  
**COMPUTER SCIENCE AND ENGINEERING**



**B.M.S. COLLEGE OF ENGINEERING**

(Autonomous Institution under VTU)

**BENGALURU-560019**

**Sep-2024 to Jan-2025**

**B.M.S. College of Engineering,**  
**Bull Temple Road, Bangalore 560019**  
(Affiliated To Visvesvaraya Technological University, Belgaum)  
**Department of Computer Science and Engineering**



**CERTIFICATE**

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Abhishek Shivanand Halagadagi (1BM22CS008)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Dr. Seema Patil Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
--	--

## Index

Sl. No.	Date	Experiment Title	Page No.
1	4-3-2025	Write a python program to import and export data using Pandas library functions	5-8
2	11-3-2025	Demonstrate various data pre-processing techniques for a given dataset	9-13
3	18-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	14-16
4	1-4-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	17-20
5	8-4-2025	Build Logistic Regression Model for a given dataset	21-23
6	15-4-2025	Build KNN Classification model for a given dataset.	24-27
7	15-4-2025	Build Support vector machine model for a given dataset	28-30
8	22-4-2025	Implement Random forest ensemble method on a given dataset.	31-32
9	22-4-2025	Implement Boosting ensemble method on a given dataset.	33-35
10	29-4-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	36-39
11	29-4-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	40-44

Github Link: <https://github.com/Abhi-008-sh/MLGLab>

## Program 1

Write a python program to import and export data using Pandas library functions

Screenshot

The image shows two pages of handwritten code and output for a Python program using Pandas. The code demonstrates how to create a DataFrame from a dictionary, import data from a CSV file, and export data to a CSV file. The output shows the resulting DataFrame for both examples.

**Page 1 (Left):**

Lab 1

```
import pandas as pd

data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [25, 30, 35],
    'City': ['New York', 'Los Angeles', 'Chicago']
}

df = pd.DataFrame(data)
df.head()
```

Output:

Name	Age	City
Alice	25	New York

→ importing dataset from sklearn dataset

```
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)

df['target'] = iris.target
print("sample data:")
df.head()
```

Output:

sepal length	sepal width	petal length	petal width
0	5.1	3.5	1.4

**Page 2 (Right):**

→ importing data from specified csv file

```
file_path = 'data.csv'
df = pd.read_csv(file_path)
print("sample data")
```

→ Document dataset

```
df = pd.read_csv('mobile-dataset-2025.csv')
print("sample data")
df.head()
```

sample data:

comp_name	model_name	weight	RAM
apple	iphone 16	174g	6GB

front.cam Back.cam processor Battery cap  
12 MP 48 MP A17 Bionic 3600mAh

size launched price (India) launched year  
6.1 inches 79999 ₹ 2024

→ Exporting data

```
1) df = pd.read_csv('sample_sales_data.csv')
2) df.to_csv('output.csv', index=False)
print("Data saved to output.csv")
```

data analysis of sales dataset

```
1) sales_df = pd.read_csv('sales_data.csv')
```



Date : \_\_\_\_\_

Page No. : \_\_\_\_\_

## 2. Summarize sales by region

```
Sales-by-re = Sales-df.groupby('Region').Sales  
.sum()
```

Sales-by-re

Region

east 770

North 16400

South 3070

west 650

## 3. Grouping by product & calculating total quantity sold

```
best-selling-pr = Sales-df.groupby('Product')  
['Quantity'].sum().sort_values(ascending  
= False)
```

best-selling-price

Product

mouse 24

laptop 17

keyboard 16

monitor 15

## ④ Saving data to a csv file

```
Sales-by-re.to_csv('Sales-by-re.csv')
```

```
best-selling-re.to_csv('best-selling-re.csv')
```

```
print("In Results saved to csv file")
```

results saved to csv file

Code:

```
import yfinance as yf

import pandas as pd

import matplotlib.pyplot as plt

tickers = ["HDFCBANK.NS", "ICICIBANK.NS", "KOTAKBANK.NS"]

data = yf.download(tickers, start="2024-01-01", end="2024-12-30", group_by='ticker')

print("First 5 rows of the dataset:")

print(data.head())

print("\nShape of the dataset:")

print(data.shape)

print("\nColumn names:")

print(data.columns)

hdfc_data = data['HDFCBANK.NS']

print("\nSummary statistics for HDFC Bank:")

print(hdfc_data.describe())

hdfc_data['Daily Return'] = hdfc_data['Close'].pct_change()

icici_data = data['ICICIBANK.NS']

print("\nSummary statistics for ICICI Bank:")

print(icici_data.describe())

icici_data['Daily Return'] = icici_data['Close'].pct_change()

kotak_data = data['KOTAKBANK.NS']

print("\nSummary statistics for Kotak Mahindra Bank:")

print(kotak_data.describe())
```

```

kotak_data['Daily Return'] = kotak_data['Close'].pct_change()

plt.figure(figsize=(14, 10))

plt.subplot(3, 2, 1)

hdfc_data['Close'].plot(title="HDFC Bank - Closing Price")

plt.subplot(3, 2, 2)

hdfc_data['Daily Return'].plot(title="HDFC Bank - Daily Returns", color='orange')

plt.subplot(3, 2, 3)

icici_data['Close'].plot(title="ICICI Bank - Closing Price")

plt.subplot(3, 2, 4)

icici_data['Daily Return'].plot(title="ICICI Bank - Daily Returns", color='orange')

plt.subplot(3, 2, 5)

kotak_data['Close'].plot(title="Kotak Mahindra Bank - Closing Price")

plt.subplot(3, 2, 6)

kotak_data['Daily Return'].plot(title="Kotak Mahindra Bank - Daily Returns", color='orange')

plt.tight_layout()

plt.show()


hdfc_data.to_csv('hdfc_bank_data.csv')

icici_data.to_csv('icici_bank_data.csv')

kotak_data.to_csv('kotak_bank_data.csv')


print("\nHDFC Bank data saved to 'hdfc_bank_data.csv'.")

print("ICICI Bank data saved to 'icici_bank_data.csv'.")

print("Kotak Bank data saved to 'kotak_bank_data.csv'.")

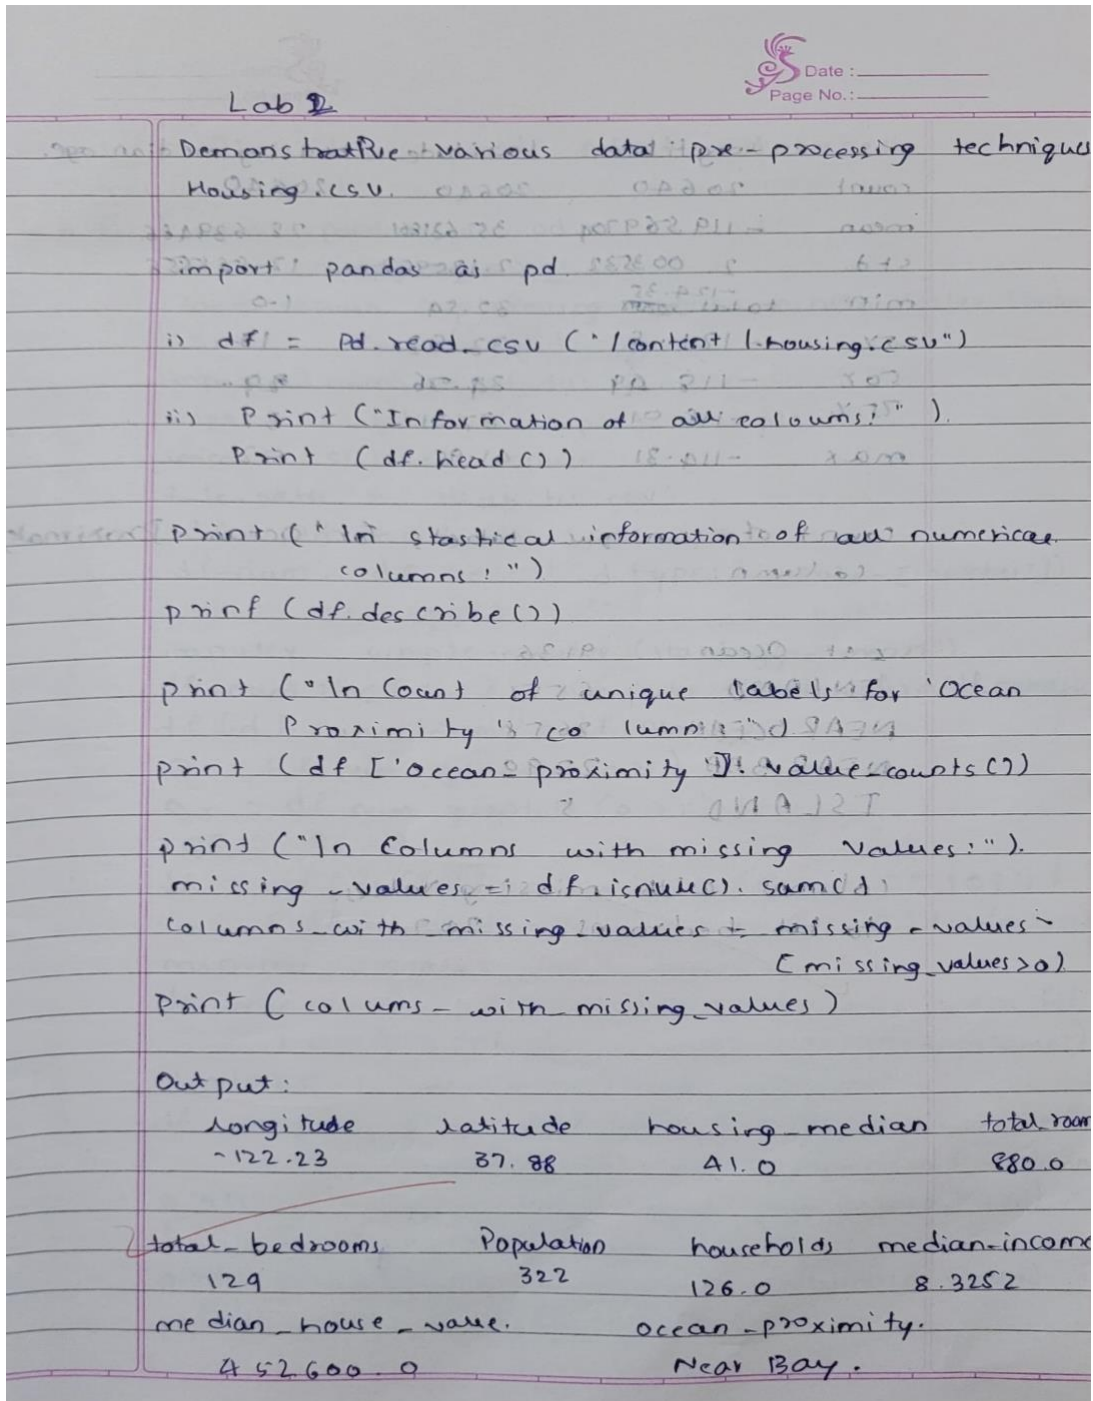
```



## Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot



The screenshot shows a handwritten program in Python for data pre-processing, titled 'Lab 2'. The program uses the pandas library to load a 'Housing.csv' file. It includes several steps: loading the data, printing information about all columns, printing the first few rows, printing statistical information for numerical columns, printing the count of unique labels for 'Ocean Proximity', printing the value counts for 'Ocean Proximity', identifying columns with missing values, and finally printing the output in a structured format.

```
Lab 2
```

```
import pandas as pd

df = pd.read_csv('content/Housing.csv')

print("Information of all columns:")
print(df.head())

print("In statistical information of all numerical columns:")
print(df.describe())

print("In count of unique labels for 'Ocean Proximity'")
print(df['Ocean Proximity'].value_counts())

print("In Columns with missing values:")
missing_values = df.isnull().sum()
columns_with_missing_values = missing_values[missing_values > 0]
print(columns_with_missing_values)
```

Output:

longitude	latitude	housing_median	total_rooms
-122.23	37.88	41.0	880.0

total_bedrooms	Population	households	median_income
129	322	126.0	8.3252

median_house_value	ocean_proximity
452600.0	Near Bay.



Date: \_\_\_\_\_  
Page No.: \_\_\_\_\_

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler, StandardScaler

df = pd.read_csv('diabetes.csv')
df_num = df.select_dtypes(include=[number])
imputer = SimpleImputer(strategy='mean')
df_num.iloc[:, :] = imputer.fit_transform(df_num)
Q1 = df_num.quantile(0.25)
Q3 = df_num.quantile(0.75)
IQR = Q3 - Q1
df = df[~((df_num < (Q1 - 1.5 * IQR)) |
          (df_num > (Q3 + 1.5 * IQR)))].astype(int)
min_max_scaler = MinMaxScaler()
df_minmax = pd.DataFrame(min_max_scaler.fit_transform(df_num), columns=df_num.columns)
standard_scaler = StandardScaler()
df_standard = pd.DataFrame(standard_scaler.fit_transform(df_minmax), columns=df_minmax.columns)
print("In processed dataset (min max scaler):")
print(df_minmax.head())
print("In processed dataset (standard scaler):")
print(df_standard.head())
```

count of unique labels for Ocean: Proximity column.

Ocean	count
INLAND	655
NEAR OCEAN	2658
NEAR BAA	2290
ISLAND	5

columns with missing values:

to tal bedrooms: 207

columns with missing values:

columns with missing values:

columns with missing values:

columns with missing values:

columns with missing values:

columns with missing values:



Date: \_\_\_\_\_  
Page No.: \_\_\_\_\_

### Diabetes and Adult datasets

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler, StandardScaler

df = pd.read_csv('diabetes.csv')
df_num = df.select_dtypes(include=[number])
imputer = SimpleImputer(strategy='mean')
df_num.iloc[:, :] = imputer.fit_transform(df_num)
Q1 = df_num.quantile(0.25)
Q3 = df_num.quantile(0.75)
IQR = Q3 - Q1
df = df[~((df_num < (Q1 - 1.5 * IQR)) |
          (df_num > (Q3 + 1.5 * IQR)))].astype(int)
min_max_scaler = MinMaxScaler()
df_minmax = pd.DataFrame(min_max_scaler.fit_transform(df_num), columns=df_num.columns)
standard_scaler = StandardScaler()
df_standard = pd.DataFrame(standard_scaler.fit_transform(df_minmax), columns=df_minmax.columns)
print("In processed dataset (min max scaler):")
print(df_minmax.head())
print("In processed dataset (standard scaler):")
print(df_standard.head())
```

④ Missing values are present in numerical columns if present, which are replaced by the mean of the respective columns.

⑤ No categorical column, no encoding

⑥ Min max scaling transform data to a fixed range (0,1) using.

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

used when: dataset does not follow a normal dist.

- Features have diff. ranges & used to be bounded, standardization transform data have zero means & unit variance

$$x' = \frac{x - \mu}{\sigma}$$

used when: dataset follows a Gaussian dist. Many ML Algos assume normality

~~Don't~~

## Adult dataset.

1) Which columns have missing values?  
workclass, occupation, native country  
Handling missing values by

→ fill with mode for categorical data  
→ drop rows if missing percentage is too high

2) Identify and encoding categorical columns are: workclass, education, marital status, occupation, relationship, race encoding strategy.

→ One-hot encoding: for nominal categorical column.  
ordinal.  
→ encoding: categorical column has a meaningful order.

③ Min Max & standardization (Difference Min Max Min Max Formula)

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

- Features have known range  
- dataset does not have normal distribution.

standardization: Formula:

$$x' = \frac{x - \mu}{\sigma}$$

\* features have unit scale.

Code:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import LabelEncoder
import seaborn as sns
import matplotlib.pyplot as plt

diabetes_data = pd.read_csv('/content/Dataset of Diabetes .csv')
adult_income_data = pd.read_csv('/content/adult.csv')

print("Diabetes Dataset:")
print(diabetes_data.head())

print("\nAdult Income Dataset:")
print(adult_income_data.head())

diabetes_numerical_cols = diabetes_data.select_dtypes(include=[np.number]).columns
diabetes_categorical_cols = diabetes_data.select_dtypes(include=[object]).columns

diabetes_imputer_num = SimpleImputer(strategy='median')
diabetes_data[diabetes_numerical_cols] =
diabetes_imputer_num.fit_transform(diabetes_data[diabetes_numerical_cols])

diabetes_imputer_cat = SimpleImputer(strategy='most_frequent')
diabetes_data[diabetes_categorical_cols] =
diabetes_imputer_cat.fit_transform(diabetes_data[diabetes_categorical_cols])

adult_income_numerical_cols = adult_income_data.select_dtypes(include=[np.number]).columns
adult_income_categorical_cols = adult_income_data.select_dtypes(include=[object]).columns

adult_income_imputer_num = SimpleImputer(strategy='median')
adult_income_data[adult_income_numerical_cols] =
adult_income_imputer_num.fit_transform(adult_income_data[adult_income_numerical_cols])

adult_income_imputer_cat = SimpleImputer(strategy='most_frequent')
adult_income_data[adult_income_categorical_cols] =
adult_income_imputer_cat.fit_transform(adult_income_data[adult_income_categorical_cols])

categorical_columns_adult = adult_income_data.select_dtypes(include=['object']).columns
label_encoder = LabelEncoder()

for col in categorical_columns_adult:
    adult_income_data[col] = label_encoder.fit_transform(adult_income_data[col])
```

```

def detect_and_remove_outliers(df):
    numerical_df = df.select_dtypes(include=[np.number])
    Q1 = numerical_df.quantile(0.25)
    Q3 = numerical_df.quantile(0.75)
    IQR = Q3 - Q1
    return df[~((numerical_df < (Q1 - 1.5 * IQR)) | (numerical_df > (Q3 + 1.5 * IQR))).any(axis=1)]

diabetes_data_cleaned = detect_and_remove_outliers(diabetes_data)
adult_income_data_cleaned = detect_and_remove_outliers(adult_income_data)

min_max_scaler = MinMaxScaler()

diabetes_numerical_cols = diabetes_data_cleaned.select_dtypes(include=[np.number]).columns
diabetes_data_normalized = diabetes_data_cleaned.copy()

diabetes_data_normalized[diabetes_numerical_cols] =
min_max_scaler.fit_transform(diabetes_data_cleaned[diabetes_numerical_cols])

adult_income_numerical_cols =
adult_income_data_cleaned.select_dtypes(include=[np.number]).columns
adult_income_data_normalized = adult_income_data_cleaned.copy()

adult_income_data_normalized[adult_income_numerical_cols] =
min_max_scaler.fit_transform(adult_income_data_cleaned[adult_income_numerical_cols])

standard_scaler = StandardScaler()

diabetes_data_standardized = diabetes_data_cleaned.copy()
diabetes_data_standardized[diabetes_numerical_cols] =
standard_scaler.fit_transform(diabetes_data_cleaned[diabetes_numerical_cols])

adult_income_data_standardized = adult_income_data_cleaned.copy()
adult_income_data_standardized[adult_income_numerical_cols] =
standard_scaler.fit_transform(adult_income_data_cleaned[adult_income_numerical_cols])

```



### Program 3

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot:

Date: \_\_\_\_\_  
Page No.: \_\_\_\_\_

Lab - 5

instance	A <sub>1</sub>	A <sub>2</sub>	classification
1	Hot	high	No
2	Hot	high	No
6	Cool	high	No
7	Hot	high	No
8	Hot	Normal	Yes

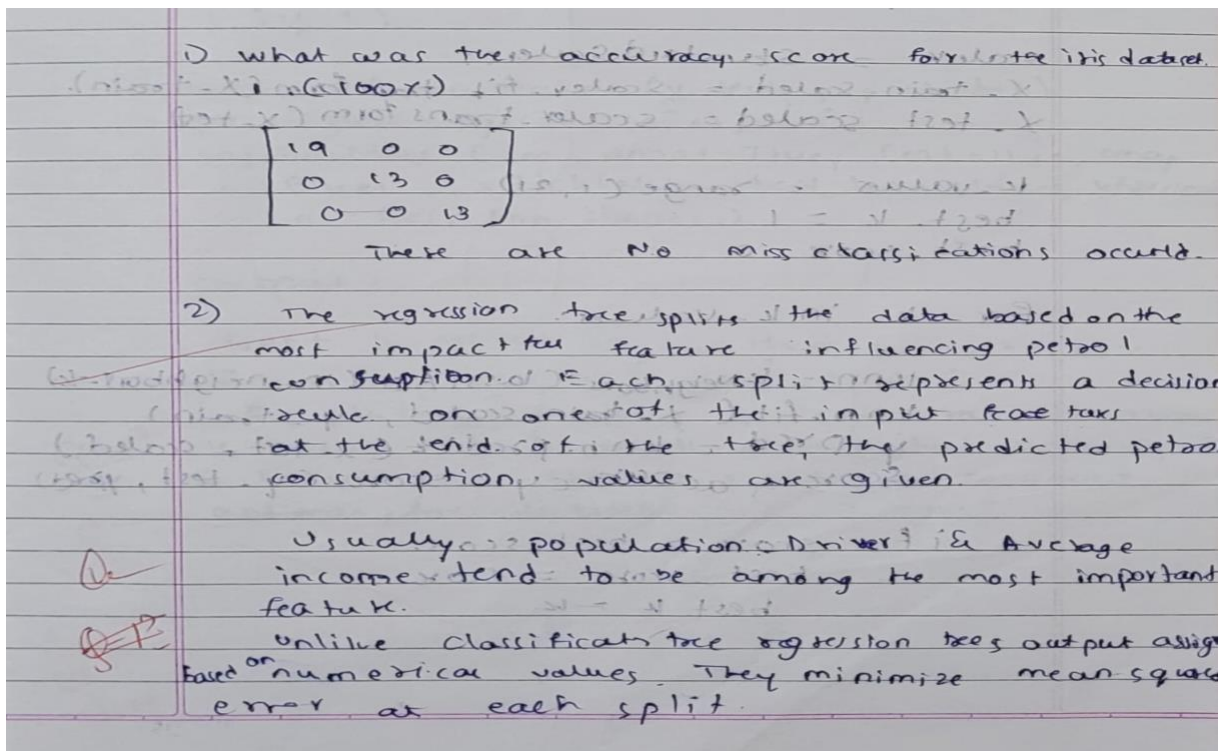
Entropy =  $-\frac{4}{5} \log \frac{4}{5} - \frac{1}{5} \log \frac{1}{5}$   
 $= 0.7219$

For A<sub>1</sub>,  
 $S_{\text{hot}}(1+, 3-) = -\frac{1}{4} \log \frac{1}{4} - \frac{3}{4} \log \frac{3}{4}$   
 $= 0.8113$   
 $S_{\text{cool}}(0+, 1-) = 0$   
 $\text{Gain}(S, A_1) = 0.7219 - \frac{4}{5} \times 0.8113 = 0.07286$

For A<sub>2</sub>,  
 $S_{\text{high}}(0+, 4-) = 0$   
 $S_{\text{normal}}(1, 0-) = 0$   
 $\text{Gain}(S, A_2) = 0.7219$

∴ A<sub>2</sub> has highest gain value it is taken as root

```
graph TD
    A2((A2)) -- high --> B1[No]
    A2 -- normal --> B2[Yes]
    B1 --- C1[1, 2, 6]
    B2 --- C2[8]
```



Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report,
mean_absolute_error, mean_squared_error
from sklearn.preprocessing import LabelEncoder

iris = pd.read_csv("/content/iris (4).csv")
drug = pd.read_csv("/content/drug.csv")
petrol = pd.read_csv("/content/petrol_consumption.csv")

X_iris = iris.iloc[:, :-1]
y_iris = iris.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X_iris, y_iris, test_size=0.2, random_state=42)

dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)
y_pred = dtc.predict(X_test)

print("Decision Tree Classification for IRIS Dataset:")
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```

print("Classification Report:\n", classification_report(y_test, y_pred))

X_drug = drug.iloc[:, :-1]
y_drug = drug.iloc[:, -1]

le = LabelEncoder()

for col in X_drug.select_dtypes(include=['object']).columns:
    X_drug[col] = le.fit_transform(X_drug[col])

X_train, X_test, y_train, y_test = train_test_split(X_drug, y_drug, test_size=0.2, random_state=42)

dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)
y_pred = dtc.predict(X_test)

print("\nDecision Tree Classification for Drug Dataset:")
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

X_petrol = petrol.iloc[:, :-1]
y_petrol = petrol.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X_petrol, y_petrol, test_size=0.2, random_state=42)

dtr = DecisionTreeRegressor()
dtr.fit(X_train, y_train)
y_pred = dtr.predict(X_test)

print("\nDecision Tree Regression for Petrol Consumption:")
print("Mean Absolute Error:", mean_absolute_error(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error:", np.sqrt(mean_squared_error(y_test, y_pred)))

```



#### Program 4

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot

Date: \_\_\_\_\_  
Page No.: \_\_\_\_\_

Lab 8 A.

$x_i$ (Weeks)	$y_i$ (Sales in thousands)
1	2
2	4
3	5
4	9

$$\beta = (X^T X)^{-1} X^T Y$$
$$X = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix} \quad Y = \begin{bmatrix} 2 \\ 4 \\ 5 \\ 9 \end{bmatrix}$$
$$X^T X = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \end{bmatrix} \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix}$$
$$X^T X = \begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix}$$
$$(X^T X)^{-1} = \begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix}^{-1} = \begin{bmatrix} 1.5 & -0.5 \\ -0.5 & 0.2 \end{bmatrix}$$
$$((X^T X)^{-1} X^T) = \begin{pmatrix} 1 & 0.5 & 0 & -0.5 \\ -0.3 & -0.1 & 0.1 & 0.3 \end{pmatrix}$$



Date : \_\_\_\_\_

Page No.: \_\_\_\_\_

$$((x^T x)^{-1} x^T) Y = \begin{bmatrix} 1 & 0.5 & 0 & -0.5 \\ -0.3 & -0.1 & 0.1 & 0.3 \end{bmatrix} \begin{bmatrix} 2 \\ 4 \\ 5 \\ 9 \end{bmatrix}$$

$$\beta = \begin{bmatrix} -0.5 \\ 2.2 \end{bmatrix} \begin{matrix} \text{intercept} \\ \text{slope} \end{matrix}$$

$$y = \beta_0 + \beta_1 x + \epsilon \quad \text{At } x = 5$$
$$y = -0.5 + 2.2x$$
$$y = -0.5 + 2.2(5)$$
$$y = 10.5$$

Ans  
11-3-25

Code:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_absolute_error
import matplotlib.pyplot as plt

hiring_data = pd.read_csv('hiring.csv')
print(hiring_data.head())
hiring_data = hiring_data.dropna()

experience_mapping = {
    'one': 1, 'two': 2, 'three': 3, 'four': 4, 'five': 5, 'six': 6, 'seven': 7, 'eight': 8,
    'nine': 9, 'ten': 10, 'eleven': 11, 'twelve': 12, 'thirteen': 13, 'fourteen': 14,
}

hiring_data['experience'] = hiring_data['experience'].replace(experience_mapping)
hiring_data['experience'] = pd.to_numeric(hiring_data['experience'], errors='coerce')

if hiring_data['experience'].isnull().any():
    print("Warning: There are still non-numeric values in the 'experience' column.")
    hiring_data = hiring_data.dropna(subset=['experience'])

X_hiring = hiring_data[['experience', 'test_score(out of 10)', 'interview_score(out of 10)']]
y_hiring = hiring_data['salary($)']

X_train_hiring, X_test_hiring, y_train_hiring, y_test_hiring = train_test_split(X_hiring, y_hiring,
test_size=0.2, random_state=42)

regressor_hiring = LinearRegression()
regressor_hiring.fit(X_train_hiring, y_train_hiring)

candidate_1 = np.array([[2, 9, 6]])
candidate_2 = np.array([[12, 10, 10]])

salary_1 = regressor_hiring.predict(candidate_1)
salary_2 = regressor_hiring.predict(candidate_2)

print(f"Predicted salary for candidate 1 (2 yr experience, 9 test score, 6 interview score):
{salary_1[0]}")
print(f"Predicted salary for candidate 2 (12 yr experience, 10 test score, 10 interview score):
{salary_2[0]}")
```

```

companies_data = pd.read_csv('/content/1000_Companies.csv')
print(companies_data.head())
companies_data = companies_data.dropna()

label_encoder = LabelEncoder()
companies_data['State'] = label_encoder.fit_transform(companies_data['State'])

X_companies = companies_data[['R&D Spend', 'Administration', 'Marketing Spend', 'State']]
y_companies = companies_data['Profit']

X_train_companies, X_test_companies, y_train_companies, y_test_companies =
train_test_split(X_companies, y_companies, test_size=0.2, random_state=42)

regressor_companies = LinearRegression()
regressor_companies.fit(X_train_companies, y_train_companies)

input_data = np.array([[91694.48, 515841.3, 11931.24, label_encoder.transform(['Florida'])[0]]])
predicted_profit = regressor_companies.predict(input_data)

print(f'Predicted profit for the given inputs (Florida State): {predicted_profit[0]}')

y_pred_hiring = regressor_hiring.predict(X_test_hiring)
mae_hiring = mean_absolute_error(y_test_hiring, y_pred_hiring)
print(f'Mean Absolute Error for Salary Prediction: {mae_hiring}')

y_pred_companies = regressor_companies.predict(X_test_companies)
mae_companies = mean_absolute_error(y_test_companies, y_pred_companies)
print(f'Mean Absolute Error for Profit Prediction: {mae_companies}')

```

## Program 5

Build Logistic Regression Model for a given dataset

Screenshot

Date: 18/03/2025  
Page No.: \_\_\_\_\_

Lab - 3

i)  $a_0 = -5$   
 $a_1 = 0.8$

ii)  $f(x) = \frac{1}{1 + \exp(-(-5 + 0.8x))}$

iii)  $f(7) = \frac{1}{1 + \exp(-(-5 + 0.8 \times 7))}$   
 $= 0.6457$

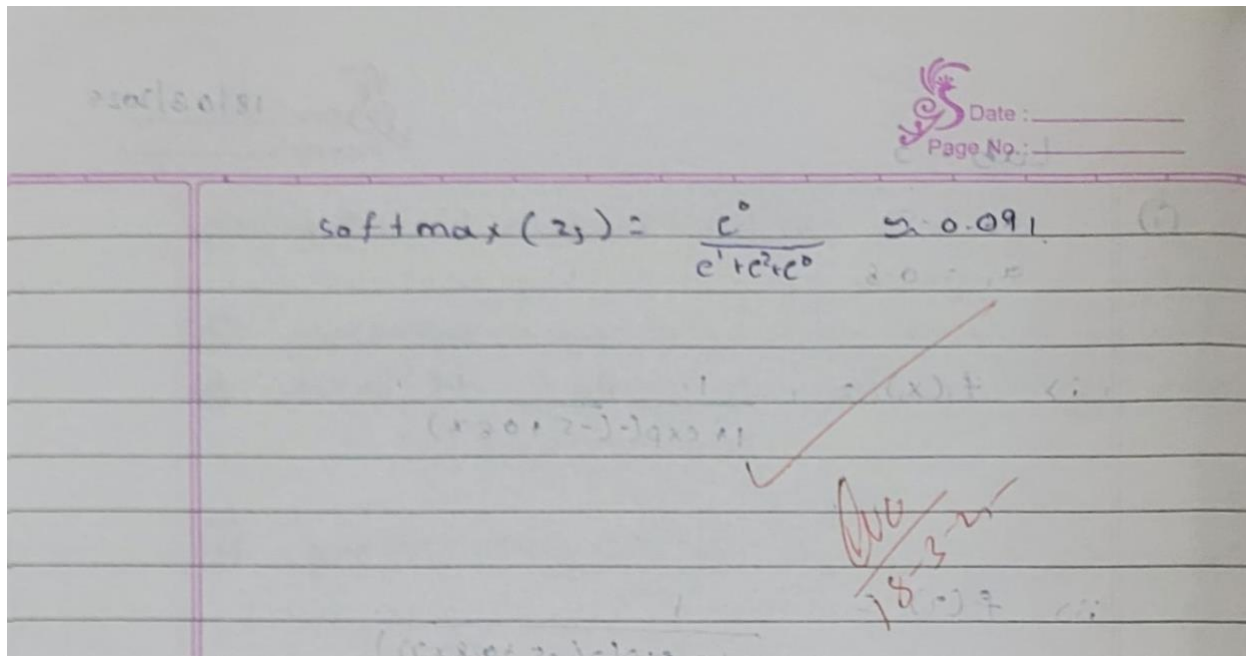
if  $f(x) < 0.5$  then the student is pass fail  
if  $f(x) \geq 0.5$  then the student is pass.

ii)  $z = [2, 1, 0]$  base  $b_0 = 2$   
 $b_1 = 1$   
 $b_2 = 0$

$\text{softmax}(z_u) = \frac{e^{z_u}}{\sum_{j=1}^n e^{z_j}}$

$\text{softmax}(z_1) = \frac{e^{z_1}}{e^{z_1} + e^{z_2} + e^{z_3}} = \frac{e^2}{e^2 + e^1 + e^0} = 0.6652$

so  $\text{softmax}(z_2) = \frac{e^{z_2}}{e^{z_1} + e^{z_2} + e^{z_3}} = \frac{e^1}{e^2 + e^1 + e^0} = 0.245$



Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
file_path = 'HR_comma_sep.csv'
data = pd.read_csv(file_path)
```

```
print(data.info())
```

```
print(data.head())
```

```
print(data.describe())
```

```
plt.figure(figsize=(8, 5))
sns.countplot(x='salary', hue='left', data=data)
plt.title('Impact of Salary on Employee Retention')
plt.xlabel('Salary')
plt.ylabel('Count')
plt.legend(title='Employee Retention', labels=['Stayed', 'Left'])
plt.show()
```

```
plt.figure(figsize=(10, 6))
```

```

sns.countplot(x='Department', hue='left', data=data)
plt.title('Impact of Department on Employee Retention')
plt.xlabel('Department')
plt.ylabel('Count')
plt.legend(title='Employee Retention', labels=['Stayed', 'Left'])
plt.xticks(rotation=45)
plt.show()

data_encoded = pd.get_dummies(data, columns=['salary', 'Department'], drop_first=True)

print(data_encoded.info())

X = data_encoded.drop('left', axis=1)
y = data_encoded['left']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

logreg = LogisticRegression(max_iter=1000)

logreg.fit(X_train_scaled, y_train)

y_pred = logreg.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy of the Logistic Regression Model: {accuracy * 100:.2f}%')

cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False, xticklabels=['Stayed', 'Left'],
yticklabels=['Stayed', 'Left'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```



## Program 6

Build KNN Classification model for a given dataset.

Screenshot

Date : \_\_\_\_\_  
Page No. : \_\_\_\_\_

Lab 6

Person	Age	Salary	Target
A	18	50	N
B	23	55	N
C	24	70	N
D	41	60	Y
E	43	70	Y
F	38	40	Y
X	35	100	?

$(x_2, y_2) = (35, 100)$

for  $(18, 50)$  dist  $= \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$   
 $= \sqrt{(35 - 18)^2 + (100 - 50)^2}$   
 $= 52.81$

for  $(23, 55)$   $= \sqrt{(35 - 23)^2 + (100 - 55)^2}$   
 $= 46.57$

for  $(24, 70)$   $= \sqrt{(35 - 24)^2 + (100 - 70)^2}$   
 $= 31.95$

for  $(41, 60)$   $= \sqrt{(35 - 41)^2 + (100 - 60)^2}$   
 $= 40.44$

for  $(43, 70)$   $= \sqrt{(35 - 43)^2 + (100 - 70)^2}$   
 $= 31.04$

for  $(38, 40)$   $= \sqrt{(35 - 38)^2 + (100 - 40)^2}$   
 $= 60.07$





Date : \_\_\_\_\_

Page No. : \_\_\_\_\_

Person	Age	Salary	Target	distance	Rank
A	18	50	N	52.81	5
B	23	55	N	46.57	4
C	24	70	N	31.95	2
D	41	60	Y	40.44	3
E	43	70	Y	31.04	1
F	38	40	Y	60.07	6
X	35	100	Y	7	

$$k=3$$

$$2-Y$$

$$1-N$$

$$k=1$$

$$-Y$$

$$\sqrt{(18-00)^2 + (50-28)^2} = \sqrt{(00, 28)} \text{ not}$$

$$k=3$$

$$-Y$$

$$\sqrt{(07-00)^2 + (81-28)^2}$$

$$\text{So, } X(35, 100) \text{ is } Y.$$

$$\sqrt{(22-00)^2 + (68-28)^2} = \sqrt{(22, 68)} \text{ not}$$

$$12.24$$

$$\sqrt{(00-00)^2 + (45-28)^2} = \sqrt{(00, 45)} \text{ not}$$

$$\sqrt{(00-00)^2 + (18-28)^2} = \sqrt{(00, 18)} \text{ not}$$

$$10.02$$

$$\sqrt{(00-00)^2 + (68-28)^2} = \sqrt{(00, 68)} \text{ not}$$

$$10.18$$

$$\sqrt{(00-00)^2 + (81-28)^2} = \sqrt{(00, 81)} \text{ not}$$

Code:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns

iris_df = pd.read_csv('/content/iris (3).csv')

print(iris_df.head())

X_iris = iris_df.drop(columns=['species'])
y_iris = iris_df['species']

X_train_iris, X_test_iris, y_train_iris, y_test_iris = train_test_split(X_iris, y_iris, test_size=0.2,
random_state=42)

scaler = StandardScaler()
X_train_iris = scaler.fit_transform(X_train_iris)
X_test_iris = scaler.transform(X_test_iris)

knn_iris = KNeighborsClassifier(n_neighbors=3)

knn_iris.fit(X_train_iris, y_train_iris)

y_pred_iris = knn_iris.predict(X_test_iris)

accuracy_iris = accuracy_score(y_test_iris, y_pred_iris)
print(f'Accuracy on Iris test data: {accuracy_iris * 100:.2f}%')

cm_iris = confusion_matrix(y_test_iris, y_pred_iris)
sns.heatmap(cm_iris, annot=True, fmt="d", cmap="Blues", xticklabels=knn_iris.classes_,
yticklabels=knn_iris.classes_)
plt.title("Confusion Matrix for Iris Dataset")
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

print("Classification Report for Iris Dataset:")
print(classification_report(y_test_iris, y_pred_iris))

diabetes_df = pd.read_csv('diabetes.csv')
print(diabetes_df.head())
```

```

X_diabetes = diabetes_df.drop(columns=['Outcome'])
y_diabetes = diabetes_df['Outcome']

X_train_diabetes, X_test_diabetes, y_train_diabetes, y_test_diabetes = train_test_split(X_diabetes,
y_diabetes, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train_diabetes = scaler.fit_transform(X_train_diabetes)
X_test_diabetes = scaler.transform(X_test_diabetes)

knn_diabetes = KNeighborsClassifier(n_neighbors=5)

knn_diabetes.fit(X_train_diabetes, y_train_diabetes)

y_pred_diabetes = knn_diabetes.predict(X_test_diabetes)

accuracy_diabetes = accuracy_score(y_test_diabetes, y_pred_diabetes)
print(f"Accuracy on Diabetes test data: {accuracy_diabetes * 100:.2f}%")

cm_diabetes = confusion_matrix(y_test_diabetes, y_pred_diabetes)
sns.heatmap(cm_diabetes, annot=True, fmt="d", cmap="Blues", xticklabels=knn_diabetes.classes_,
yticklabels=knn_diabetes.classes_)
plt.title("Confusion Matrix for Diabetes Dataset")
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

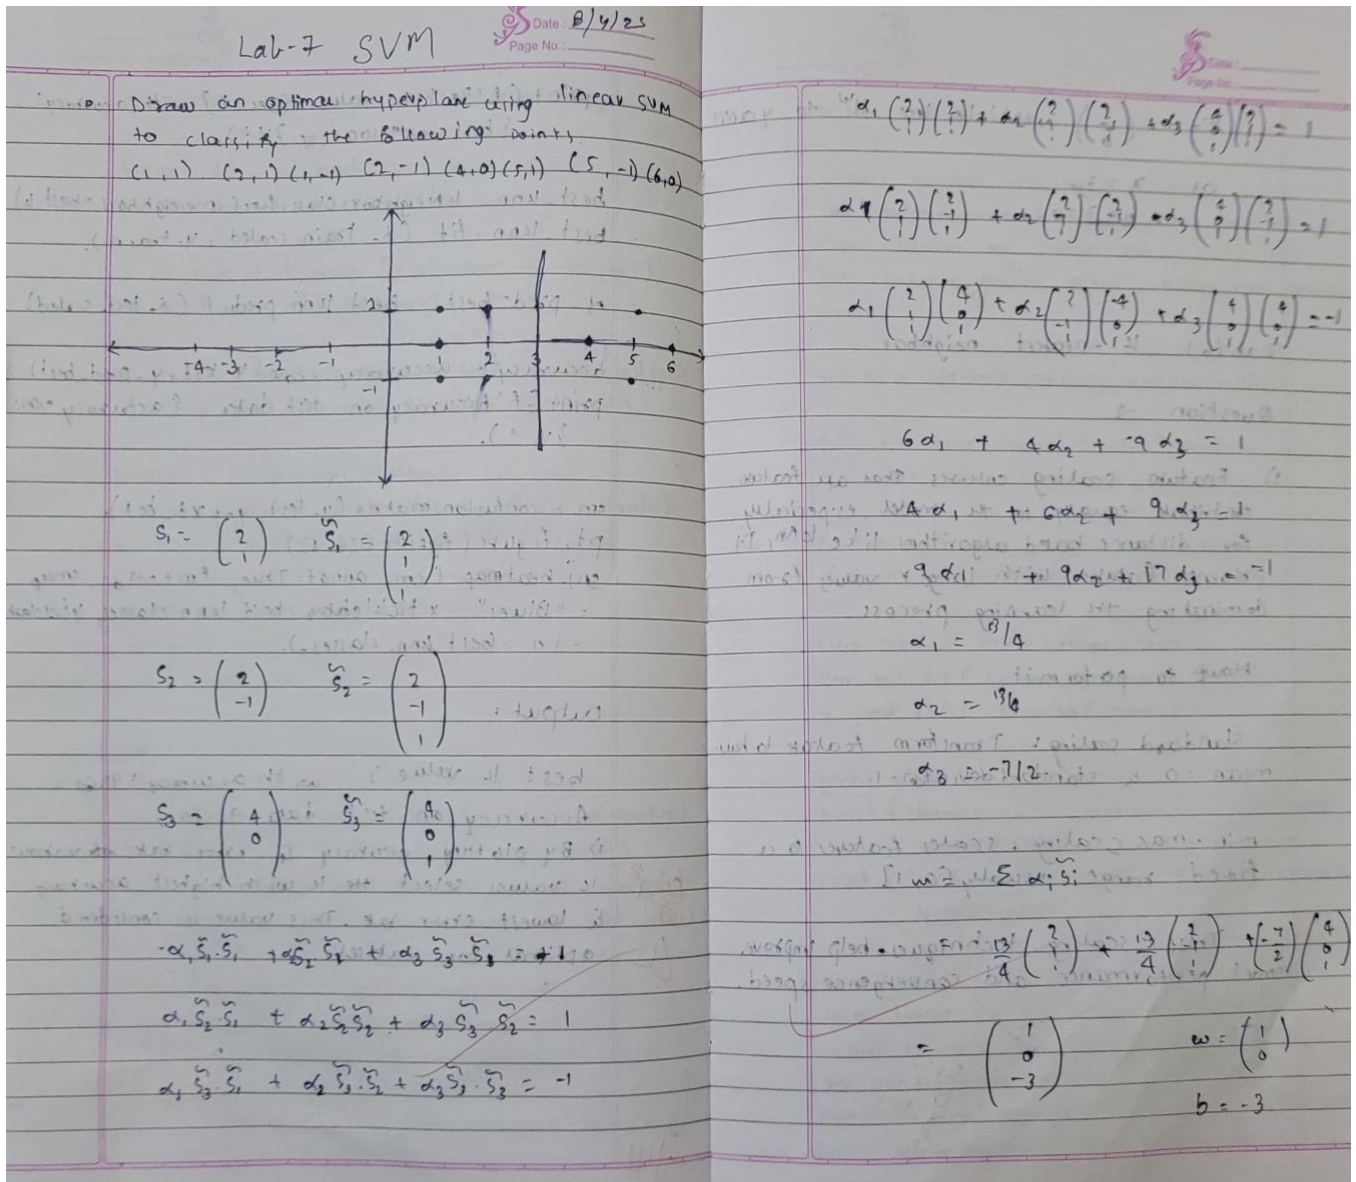
print("Classification Report for Diabetes Dataset:")
print(classification_report(y_test_diabetes, y_pred_diabetes))

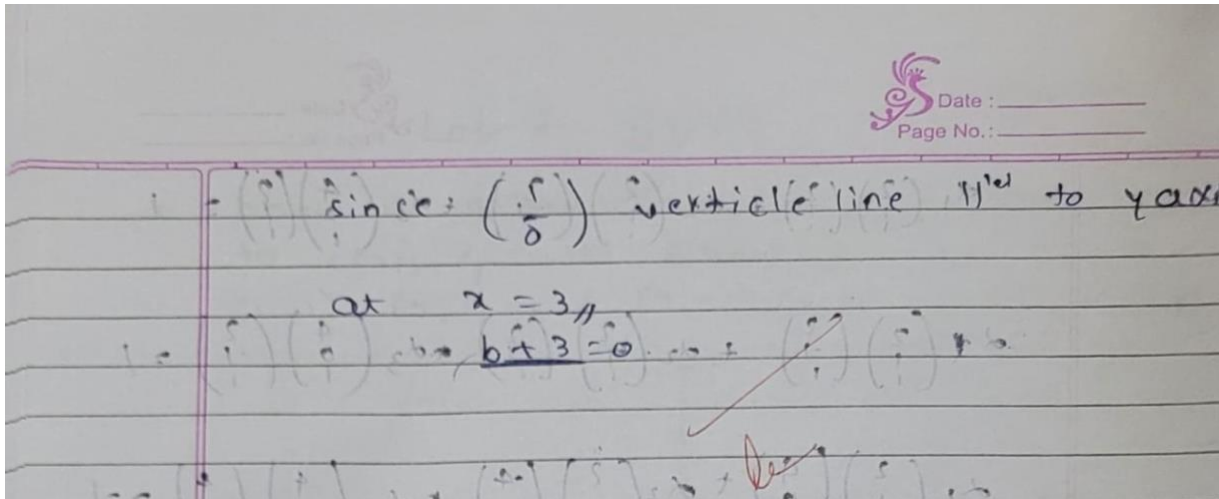
```

## Program 7

Build Support vector machine model for a given dataset

Screenshot





Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_score, roc_curve
from sklearn.preprocessing import LabelEncoder, label_binarize
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

df = pd.read_csv("/content/letter-recognition.csv")

top_classes = df['letter'].value_counts().head(5).index.tolist()
df = df[df['letter'].isin(top_classes)]

X = df.iloc[:, 1:]
y = df.iloc[:, 0]

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

y_bin = label_binarize(y_encoded, classes=np.unique(y_encoded))
n_classes = y_bin.shape[1]

X_train, X_test, y_train, y_test_bin = train_test_split(X, y_bin, test_size=0.2, random_state=42)

svm_model = SVC(kernel='linear', probability=True)
svm_model.fit(X_train, y_train.argmax(axis=1))
y_score = svm_model.predict_proba(X_test)

y_pred = svm_model.predict(X_test)
```

```

y_true = y_test_bin.argmax(axis=1)

print("Accuracy:", accuracy_score(y_true, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_true, y_pred))

plt.figure()
for i in range(n_classes):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    auc = roc_auc_score(y_test_bin[:, i], y_score[:, i])
    plt.plot(fpr, tpr, label=f"{label_encoder.inverse_transform([i])[0]} AUC={auc:.2f}")
plt.plot([0, 1], [0, 1], 'k--')
plt.title("ROC Curve (Top 5 Classes)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.tight_layout()
plt.show()

macro_auc = roc_auc_score(y_test_bin, y_score, average="macro")
print("Macro AUC Score:", macro_auc)

```

## Program 8

Implement Random forest ensemble method on a given dataset.

Screenshot

Date: 15/04/2025  
Page No.: \_\_\_\_\_

### Lab - 8 (Random Forest)

- 1) Decision Tree vs Random Forest:
  - single tree vs Multiple.
  - less accurate vs More accurate
  - Fast to train vs slower to train
  - single prediction vs majority vote
- 2) Parameters of Random Forest Classification:
  - n-estimate: no. of trees in forest
  - criterion: measure to nearest quality of split
  - max-depth: min depth of tree
  - min-sample-split: min sample req to split node
  - Max feature: no. of feature to look for best fit
- 3) Algorithm:
  - 1) training dataset = split into train & test
  - 2) for n times:
    - randomly select sample with replacement
    - grow a decision tree
    - at each split choose random subset of features
    - split nodes using best feature
  - 3) Aggregate prediction:
    - classification - majority vote
    - regression - average

Q. O/p prediction.

Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn import preprocessing

df = pd.read_csv('/content/train.csv')

X = df.iloc[:, :-1]
y = df.iloc[:, -1]

for column in X.columns:
    if X[column].dtype == 'object':
        le = preprocessing.LabelEncoder()
        X[column] = le.fit_transform(X[column])

if y.dtype == 'object':
    le = preprocessing.LabelEncoder()
    y = le.fit_transform(y)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

rf_classifier = RandomForestClassifier(random_state=42)
rf_classifier.fit(X_train, y_train)

y_pred = rf_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

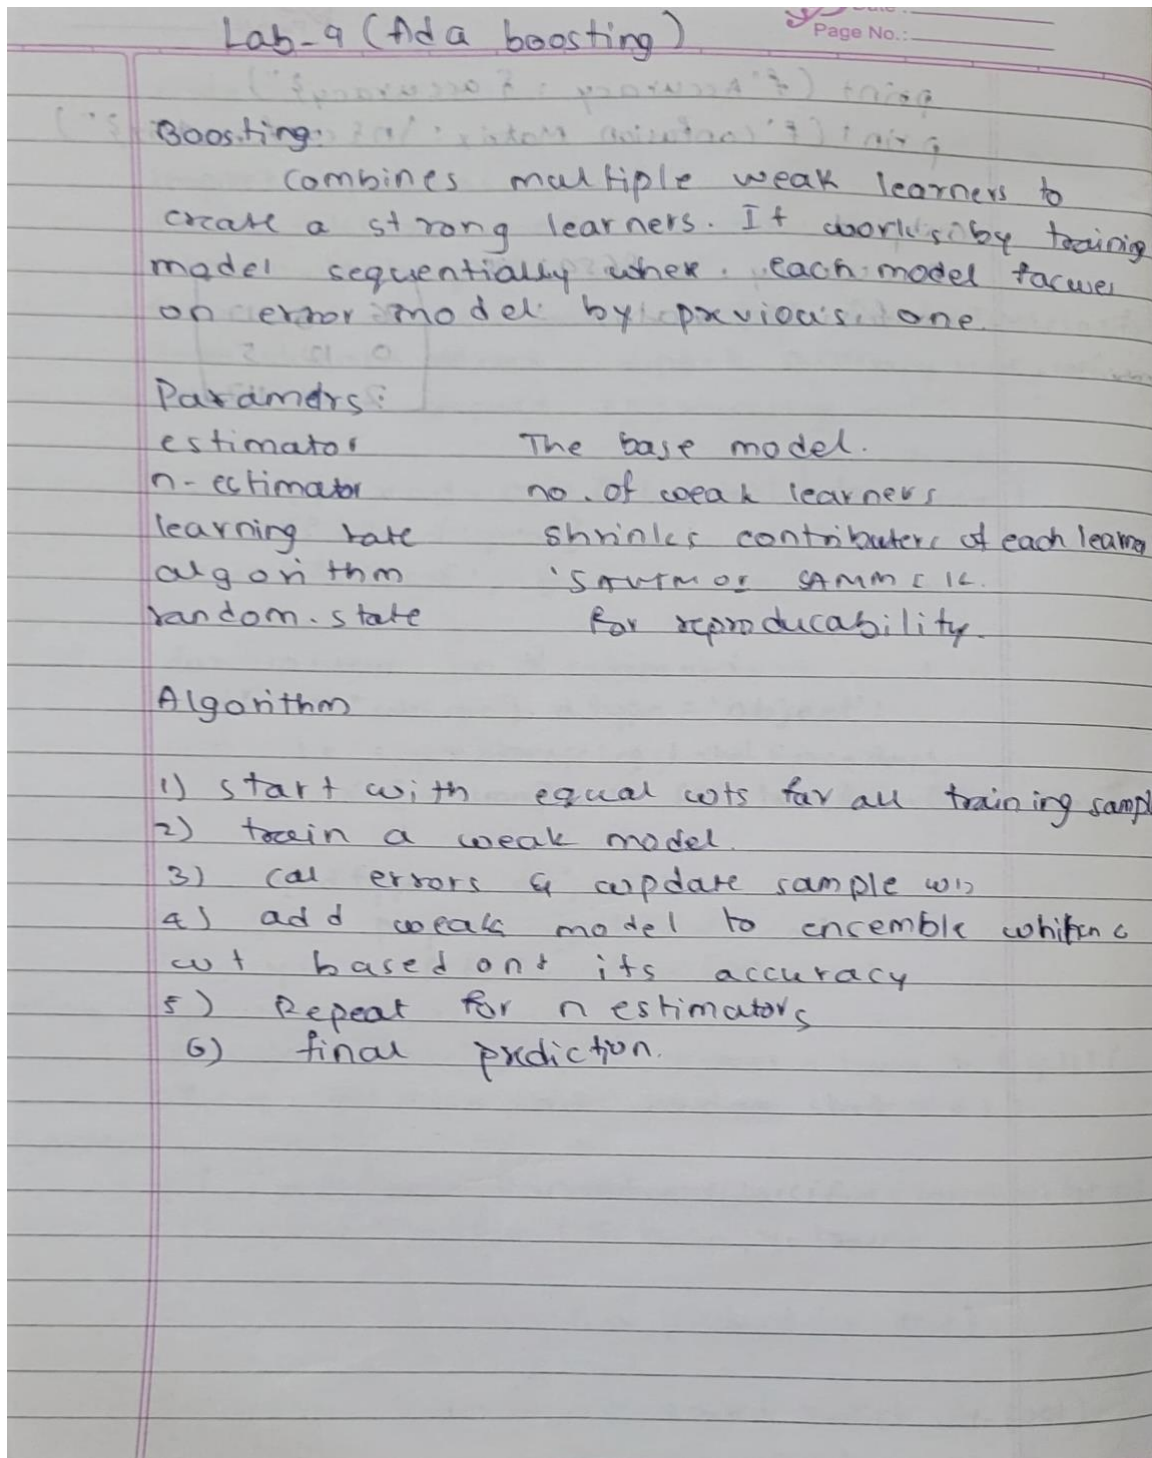
print(f'Accuracy: {accuracy}')
print(f'Confusion Matrix:\n{conf_matrix}')
```



## Program 9

Implement Boosting ensemble method on a given dataset.

Screenshot



Code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

iris = load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

results = []

n_estimators_list = [10, 50, 100]
learning_rates = [0.01, 0.1, 1]

for n in n_estimators_list:
    for lr in learning_rates:
        tree_base = DecisionTreeClassifier(max_depth=1)
        model = AdaBoostClassifier(estimator=tree_base, n_estimators=n, learning_rate=lr,
random_state=42)
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        acc = accuracy_score(y_test, y_pred)
        results.append({
            'Base': 'DecisionTree',
            'n_estimators': n,
            'learning_rate': lr,
            'Accuracy': acc
        })

for n in n_estimators_list:
    for lr in learning_rates:
        log_reg_base = LogisticRegression(max_iter=1000)
        model = AdaBoostClassifier(estimator=log_reg_base, n_estimators=n, learning_rate=lr,
random_state=42)
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        acc = accuracy_score(y_test, y_pred)
        results.append({
            'Base': 'LogisticRegression',
```

```
        'n_estimators': n,  
        'learning_rate': lr,  
        'Accuracy': acc  
    })  
  
results_df = pd.DataFrame(results)  
print(results_df)  
  
import seaborn as sns  
plt.figure(figsize=(12, 6))  
sns.barplot(x='n_estimators', y='Accuracy', hue='Base', data=results_df, ci=None)  
plt.title('AdaBoost Accuracy with Different Estimators and n_estimators')  
plt.show()
```

## Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Screenshot

**Left Page:**

```
optimal_k = 3
kmeans = KMeans(n_clusters=optimal_k, random_state=0)
kmeans.fit(X_train)
y_pred = kmeans.predict(X_test)
print(f'Predicted clusters for test data: {y_pred}')

output:
Predicted clusters for Test Data: [1 0 1 0 2 0 1 1 0 0]
```

**Q1) Algorithm:**

- 1) select number  $k$  to decide the no. of clusters.
- 2) select random  $k$  points or centroids.
- 3) Assign each data points to their closest centroid, which will form the predefined  $k$  clusters.
- 4) Calculate the variance & place a new centroid of each cluster.
- 5) Repeat the third steps, which means reassign each datapoint to new closest centroid of each cluster.
- 6) If any reassignment occurs then go to step-4 else go to Finish.
- 7) The model is ready.

**Q2) How to determine no. of clusters?**

- Elbow method
- Silhouette score
- Gaps statistic

**Right Page:**

**Q2) Formula for Sum of squared error**

Plot  $SSE$  vs no. of clusters.

$$SSE = \sum_{k=1}^K \sum_{i \in C_k} ||x_i - \mu_k||^2$$

The plot of  $SSE$  vs no. of clusters.

**Step 1:** Run the k-means algorithm for a range of cluster values.

- 1) calculate  $SSE$  for each no. of clusters.
- 2) Plot  $SSE$  vs no. of clusters.
- 3) calculate the variance & place a new centroid of each cluster.
- 4) Repeat the step 3, which means reassign each datapoint to the new closest centroid of each cluster.
- 5) If any reassignment occurs then go to step 3.
- 6) Finish.
- 7) The model is ready.

**Q3) Elbow Technique:**

- i) Run k means for range of values of  $k$ .
- ii) compute the within clusters sum of squares each value of  $k$ .
- iii) Plot the wss against the no. of clusters.
- iv) Look for the elbow point the value of  $k$  at which the rate of decrease in wss slowed significantly.

- Q) Discuss all parameters used in k-means++
- in-cluster: No. of clusters to form.
  - init: k-means++ random an array of shape  $(k, 2)$
  - n-init: No. of times the k-means algorithm will be run with different centroid seeds.
  - max\_iter: Max no. of iterations of the k-means algorithm.

vi) tol: Relative tolerance with regards to inertia to declare convergence.

vii) Random\_state: controls the regards to inertia to declare convergence.

viii) Algorithm: k-means implementation.

- Q) cluster 8 points (with  $(x, y)$  representing location) into 3 clusters:  $A_1(2, 10)$ ,  $A_2(7, 5)$ ,  $A_3(6, 4)$ ,  $A_4(5, 8)$ ,  $A_5(7, 5)$ ,  $A_6(6, 4)$ ,  $A_7(1, 2)$ ,  $A_8(4, 9)$   
initial clusters:  $A_1(2, 10)$ ,  $A_2(5, 8)$ ,  $A_3(1, 2)$

→ Calculation of distance b/w point  $A_1(2, 10)$  &

$$P(A_1, C_1) = |x_2 - x_1| + |y_2 - y_1| = 0$$

$$C_1(5, 8)$$

$$P(A_1, C_2) = 5$$

$$C_2(1, 2)$$

$$P(A_1, C_3) = |2 - 1| + |10 - 2| = 8 + 1 = 9$$

we calculate distance of other points from each of the center of 3 clusters.

for cluster 01

- we have only one point  $A_1(2, 10)$  in cluster
- cluster center remains the same.

iteration - 1

Given Pts	Dist from $C_1$	Dist from $C_2$	Dist from $C_3$	Point belongs to
$A_1(2, 10)$	0	5	9	$C_1$
$A_2(7, 5)$	5	6	14	$C_3$
$A_3(6, 4)$	12	7	9	$C_2$
$A_4(5, 8)$	5	0	10	$C_2$
$A_5(7, 5)$	10	5	9	$C_2$
$A_6(6, 4)$	10	5	7	$C_2$
$A_7(1, 2)$	9	10	0	$C_3$
$A_8(4, 9)$	3	2	10	$C_2$

for 02

center of cluster 02:

$$(8 + 5 + 7 + 6 + 4) / 5 = (30) / 5 = (6, 6)$$

for cluster 03:

$$(2 + 1) / 2, (5 + 2) / 2 = (1.5, 3.5)$$

this is completion of iteration 01



iteration 02:

Given Pts	Dist from $C_1$	Dist from $C_2$	Dist from $C_3$	Point Belongs to cluster
$A_1(2,10)$	0	8	7	$C_1$
$A_2(2,5)$	5	5	2	$C_3$
$A_3(8,4)$	12	4	7	$C_2$
$A_4(5,8)$	5	3	8	$C_2$
$A_5(7,5)$	10	2	7	$C_2$
$A_6(6,4)$	10	2	5	$C_2$
$A_7(1,2)$	9	9	2	$C_3$
$A_8(8,9)$	3	5	8	$C_1$

for  $C_1$  center =  $((2+9)/2, (10+9)/2) = (3.5, 9.5)$   
 for  $C_2$  center =  $((8+5+7+6)/4, (4+8+5+4)/4) = (6.5, 5.25)$   
 for  $C_3$  center =  $((2+1)/2, (5+2)/2) = (1.5, 3.5)$

iteration 03:

Given Pts	Dist from $C_1$	Dist from $C_2$	Dist from $C_3$	Point Belongs to cluster
$A_1(2,10)$	1.5	9.25	7	$C_1$
$A_2(2,5)$	5.5	4.75	2	$C_3$
$A_3(8,4)$	10.5	2.75	7	$C_2$
$A_4(5,8)$	3.5	4.75	8	$C_1$
$A_5(7,5)$	8.5	0.75	7	$C_2$
$A_6(6,4)$	8.5	1.75	5	$C_2$
$A_7(1,2)$	9.5	8.75	2	$C_3$
$A_8(8,9)$	1.5	6.25	8	$C_1$

for  $C_1$   $((2+9)/3, (10+9)/3) = (3.66, 9)$   
 for  $C_2$   $((8+5+4)/3, (4+8+4)/3) = (7.4, 3.3)$   
 for  $C_3$   $((2+1)/2, (5+2)/2) = (1.5, 3.5)$

iteration 04

Given Pts	Dist from $C_1$	Dist from $C_2$	Dist from $C_3$	Point Belongs to cluster
$A_1(2,10)$	1.66	10.07	8	$C_1$
$A_2(2,5)$	5.66	5.67	2	$C_3$
$A_3(8,4)$	9.34	1.33	7	$C_2$
$A_4(5,8)$	2.34	5.67	8	$C_1$
$A_5(7,5)$	7.34	0.67	7	$C_2$
$A_6(6,4)$	7.34	1.33	6	$C_2$
$A_7(1,2)$	9.66	8.37	2	$C_3$
$A_8(8,9)$	0.34	7.66	8	$C_1$

→ center of clusters are  $(3.66, 9)$   $(7.4, 3.3)$   $(1.5, 3.5)$

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

data = {
    'Name': [f'Person_{i+1}' for i in range(50)],
    'Age': np.random.randint(18, 70, size=50),
    'Income': np.random.randint(20000, 120000, size=50)
}

df = pd.DataFrame(data)

df.to_csv('income.csv', index=False)

df = pd.read_csv('income.csv')

X = df[['Age', 'Income']]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test = train_test_split(X_scaled, test_size=0.2, random_state=42)

sse = []
k_range = range(1, 11)
for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_train)
    sse.append(kmeans.inertia_)

plt.plot(k_range, sse, marker='o')
plt.title('SSE vs Number of Clusters')
plt.xlabel('Number of Clusters')
plt.ylabel('Sum of Squared Errors (SSE)')
plt.show()

optimal_k = 3
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
kmeans.fit(X_train)
y_pred = kmeans.predict(X_test)

print(f'Predicted Clusters for Test Data: {y_pred}')
```

## Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method.

Screenshot

Date: 6/5/2025  
Page No.: \_\_\_\_\_

Lab - 11

- 1) calculate mean
- 2) calculation of covariance matrix
- 3) Eigenvalues of the covariance matrix
- 4) Computation of the eigenvectors - Unit eigenvectors
- 5) computation of first principal component
- 6) Geometrical meaning of first principle components.

Given the data in Table, reduce the dimension (from 2 to 1) using the principle component Analysis Algorithm.

Feature	example 1	2	3	4
$x_1$	4	8	13	7
$x_2$	11	4	5	14

Step 1: calculate mean

$$\bar{x}_1 = \frac{4+8+13+7}{4} = 8$$
$$\bar{x}_2 = \frac{11+4+5+14}{4} = 8.5$$

Step 2: calculate the covariance matrix.

$$\text{cov}(x_1, x_1) = \frac{1}{N-1} \sum_{i=1}^N (x_{1i} - \bar{x}_1)^2$$
$$= \frac{1}{3} ((4-8)^2 + (8-8)^2 + (13-8)^2 + (7-8)^2)$$
$$= \frac{1}{3} (16 + 0 + 25 + 1) = \frac{42}{3} = 14$$

(11-)(11-)(11-)(11-)(11-)(11-)(11-)(11-)(11-)(11-)



$$\text{cov}(x_1, x_2) = \frac{1}{N-1} \sum_{i=1}^N (x_{1i} - \bar{x}_1)(x_{2i} - \bar{x}_2)$$

$$= \frac{1}{3} \left[ (11-8.5)^2 + (4-8)^2 + (5-8.5)^2 + (10-8.5)^2 \right]$$

$$\text{cov}(x_2, x_1) = \text{cov}(x_1, x_2)$$

$$\text{cov}(x_2, x_2) = \frac{1}{N-1} \sum_{i=1}^N (x_{2i} - \bar{x}_2)^2$$

$$= \frac{1}{3} \left[ (11-8.5)^2 + (4-8)^2 + (5-8.5)^2 + (10-8.5)^2 \right]$$

$$= 23$$

⇒ The covariance matrix is

$$S = \begin{bmatrix} \text{cov}(x_1, x_1) & \text{cov}(x_1, x_2) \\ \text{cov}(x_2, x_1) & \text{cov}(x_2, x_2) \end{bmatrix}$$

$$= \begin{bmatrix} 10 & -11 \\ -11 & 23 \end{bmatrix}$$

step 3: finding the eigenvalues

$$0 = \det(S - \lambda I)$$

$$= \begin{vmatrix} 10-\lambda & -11 \\ -11 & 23-\lambda \end{vmatrix}$$

$$= (10-\lambda)(23-\lambda) - (-11)(-11)$$

$$= \lambda^2 - 33\lambda + 20$$

solving the characteristic equation we get

$$\lambda = \frac{1}{2} (33 \pm \sqrt{565})$$

$$= 30.3849, 6.6151$$

"(h-1) + 11 = λ, λ₂" (say)

step 4: computation of the eigenvector

$$\lambda = \lambda_1$$

$$U = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = (S - \lambda_1 I)U$$

$$= \begin{bmatrix} 10-\lambda_1 & -11 \\ -11 & 23-\lambda_1 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$$

$$= \begin{bmatrix} (10-\lambda_1)u_1 - 11u_2 \\ -11u_1 + (23-\lambda_1)u_2 \end{bmatrix}$$

$$(10-\lambda_1)u_1 - 11u_2 = 0$$

$$-11u_1 + (23-\lambda_1)u_2 = 0$$

$$\frac{u_1}{11} = \frac{u_2}{10-\lambda_1} = t$$

$$u_1 = 11t, u_2 = (10-\lambda_1)t$$

taking t is any real number



Date : \_\_\_\_\_

Page No.: \_\_\_\_\_

Step 4: taking  $t = 1$  and substituting in the equation

$$(2226 + 58U_1) = \begin{bmatrix} 11 \\ 14 - d_1 \end{bmatrix}$$

Step 5: (part)  $\|U_1\| = \sqrt{11^2 + (14 - d_1)^2}$

$$\begin{aligned} &= \sqrt{11^2 + (14 - 30.3849)^2} \\ &= 19.7348 \end{aligned}$$

$$\Rightarrow e' = \begin{bmatrix} 11 / \|U_1\| \\ (14 - d_1) / \|U_1\| \end{bmatrix} = \begin{bmatrix} 0.5574 \\ -0.8303 \end{bmatrix}$$

$$e_1 = \begin{bmatrix} 0.5574 \\ -0.8303 \end{bmatrix}$$

$$e_2 = \begin{bmatrix} 0.8303 \\ 0.5574 \end{bmatrix}$$

$\Rightarrow$  by computing similar steps we get

$$\begin{bmatrix} U_1 \\ U_2 \end{bmatrix} = \begin{bmatrix} 0.5574 \\ -0.8303 \end{bmatrix}$$

Step 5: computation of first principal component

$$0 = \lambda(h - 80) + \lambda \|U_1\|$$

$$\begin{bmatrix} x_{11} \\ x_{21} \end{bmatrix}$$

$$e_1 \begin{bmatrix} x_{11} - \bar{x}_1 \\ x_{21} - \bar{x}_2 \end{bmatrix} = (0.5574 - 0.8303) \begin{bmatrix} x_{11} - \bar{x}_1 \\ x_{21} - \bar{x}_2 \end{bmatrix}$$

$$\text{and then } 10\% \text{ of } 0.5574(x_{11} - \bar{x}_1) + 0.8303(x_{21} - \bar{x}_2)$$

Code:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA
from scipy import stats

df = pd.read_csv('heart (2).csv')

z_scores = np.abs(stats.zscore(df.select_dtypes(include=[np.number])))
df_no_outliers = df[(z_scores < 3).all(axis=1)]

df_cleaned = df_no_outliers.copy()
for col in df_cleaned.select_dtypes(include='object').columns:
    df_cleaned[col] = LabelEncoder().fit_transform(df_cleaned[col])

X = df_cleaned.drop('HeartDisease', axis=1)
y = df_cleaned['HeartDisease']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42,
stratify=y)

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC()
}

print("Accuracy without PCA:")
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    print(f'{name}: {acc:.4f}')

pca = PCA(n_components=5)
X_pca = pca.fit_transform(X_scaled)
X_train_pca, X_test_pca, y_train, y_test = train_test_split(X_pca, y, test_size=0.2, random_state=42,
stratify=y)
```

```
print("\nAccuracy with PCA:")
for name, model in models.items():
    model.fit(X_train_pca, y_train)
    y_pred = model.predict(X_test_pca)
    acc = accuracy_score(y_test, y_pred)
    print(f'{name}: {acc:.4f}')
```